

PVMv2 Specification

Thomas Bytheway
Lucian Carata

September 14, 2018

1 Introduction

2 Abstract Types

The PVMv2 model defines a number of different abstract types which are used to both represent data in the resulting graph format, and as a basis for deriving concrete types to represent real trace entities. Figure 1 shows a UML representation of the different types and how they are related. Below we will go into more detail on each of the types in turn.




Figure 1: A diagram of the basic abstract types in PVMv2

Entity is the root abstract type. The common feature of all entities is that they possess a uniquely identifying name that is either totally unique for some subtypes, or is maintained for version chains for other subtypes. In general concrete types should not be derived from Entity unless there is no better option available.

Actor entities are active entities that perform actions, generally processing units such as unix processes, though they can also represent more abstract notions of processing. They are directly comparable to subjects in the orange book. They act upon objects, or other actors.

Object Passive entities that are acted upon by actors. Generally data carrying entities.

Store Object type that stores data internally. Versions from first write to last write, using EditSessions.

ChannelEP Object type that data flows through without internal storage. Does not version. May participate in connections.

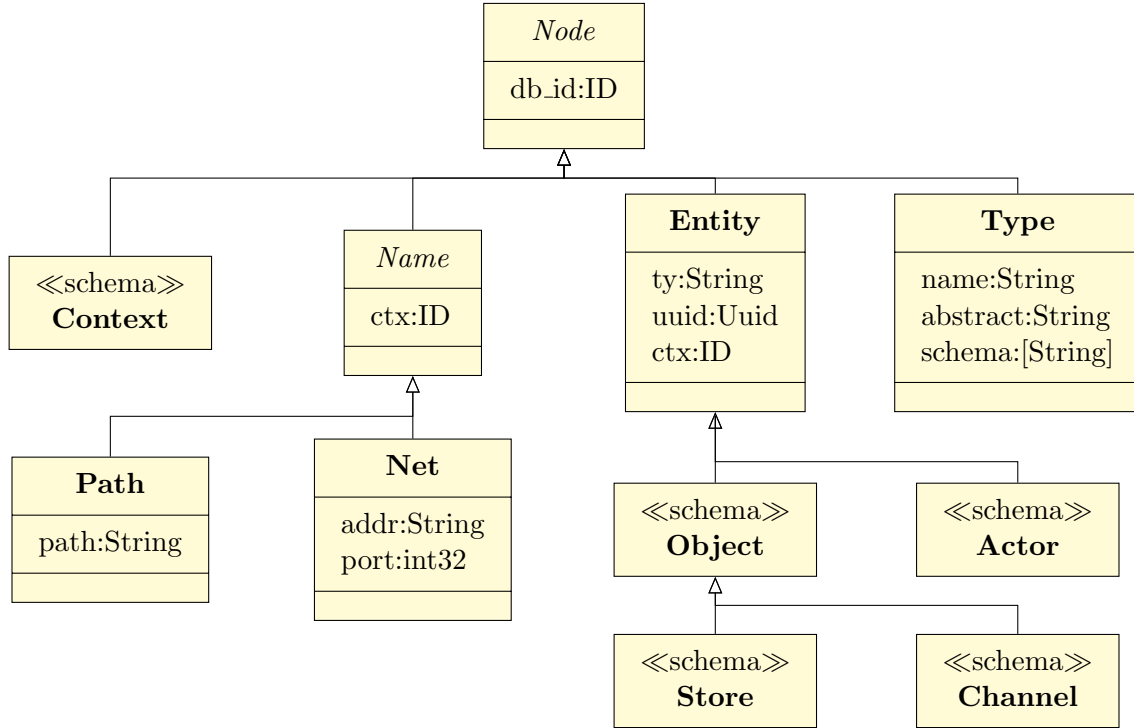


Figure 2: A diagram of the potential nodes in a PVMv2 database.

3 Graph Schema

3.1 Nodes

3.1.1 Schema Nodes

3.1.2 Context Nodes

3.1.3 Entity Nodes

3.1.4 Non-canonical Names

Path Names

Network Addresses

3.2 Relations

INF Indicates that the source has potentially imparted data or control to the destination.

NAMED Indicates that the source has been mentioned by the destination non-canonical name.

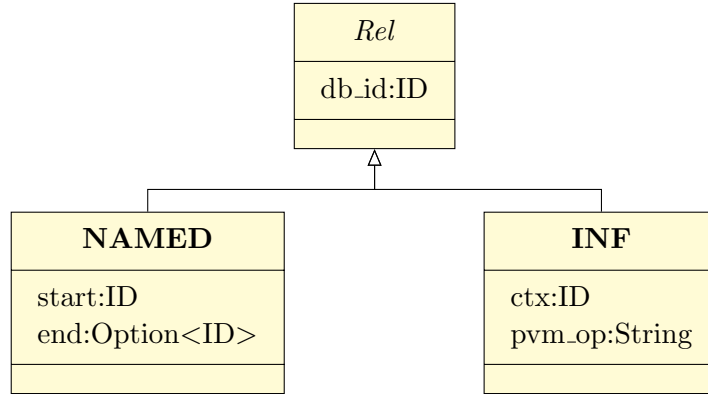


Figure 3: A diagram of the potential relations in a PVMv2 database.

4 Mappings

4.1 Concrete Type Declaration

4.2 Verbs

Declare Forces the creation of entity if it does not already exist.

Sink Called when as part of a mapped function, an actor sinks data to an entity as an atomic operation. Causes Stores to version, creating a new Store entity connected to the previous one by an INF relation. Creates an INF relation from actor to entity.

Source Called when as part of a mapped function, an actor sources data from an entity. Creates an INF relation from entity to actor.

Connect Called as part of a mapped function to declare that two channel nodes are connected to each other and that data can flow between them. The dir argument indicates the direction of flow, either mono or bi. Mono indicates that data may only flow from the first channel to the second, bi indicates that data may flow in both directions.

Mention

Unlink

Property

5 A Worked Example